

# GRID


Grupo de Investigación en Redes, Información y Distribución

## ***Programación para Unidades de Procesamiento Gráfico de Propósito General***

### ***Una oportunidad para el cálculo científico (Parte III)***


JULIAN ESTEBAN GUTIERREZ POSADA  
CHRISTIAN ANDRÉS CANDELA URIBE  
LUIS EDUARDO SEPÚLVEDA RODRÍGUEZ

Cali, agosto 2016



# Contenido

- Makefile
- Transpuesta de una matriz
- Producto de dos matrices (dos versiones)



# Makefile

make

make all

make build

make run

make clean

```
# Nombre del programa
TARGET      = Demo

EXT         = .cu


COMPILER    = nvcc

all: build run

build:
    $(COMPILER) -o $(TARGET) $(TARGET)$(EXT)

run:
    @clear
    @./$(TARGET)

clean:
    rm -f $(TARGET) $(TARGET).o *~
```



# Makefile

```
# Nombre del programa
TARGET      = SumVec-Sec


# Argumentos del programa
ARG         = 2 4 8 16 32 64 128 256 512 1024
EXT         = .cu
COMPILER    = nvcc

all: build run

build:
    $(COMPILER) -o $(TARGET) $(TARGET)$(EXT)


run:
    @clear
    @date | tee Time.txt
    @echo "\n\t N \t\t\t\t Estado \t T.Real(sg)\tT.Usuario(sg)\tT.Kernel(sg)\tCPU\tMemoria(KB)" | \
        tee -a Time.txt
    @for argument in $(ARG) ;\
        do \
            /usr/bin/time -f "\t %e \t\t %U \t\t %S \t\t %P \t\t %M" \
            ./${TARGET} $$argument 2>&1 | tee -a Time.txt ;\
        done

clean:
    rm -f $(TARGET) $(TARGET).o *~
```

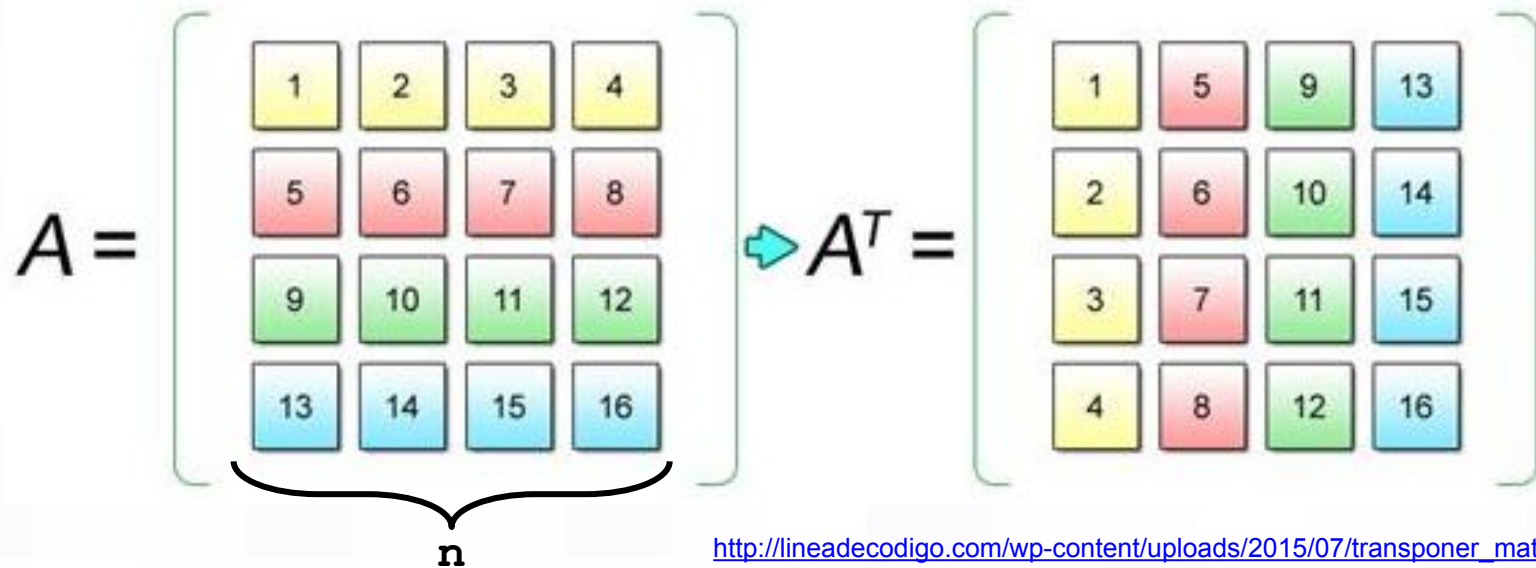


# Makefile - Salida

N	Estado	T.Real(sg)	T.Usuario(sg)	T.Kernel(sg)	CPU	Memoria(KB)
2	Ok	0.36	0.01	0.05	18%	25584
4	Ok	0.05	0.00	0.04	94%	25700
8	Ok	0.04	0.00	0.03	100%	25656
16	Ok	0.03	0.00	0.02	97%	25728
32	Ok	0.03	0.00	0.02	91%	25820
64	Ok	0.03	0.00	0.02	91%	23652
128	Ok	0.07	0.03	0.03	97%	27916
256	Ok	0.21	0.13	0.06	97%	26064
512	Ok	0.99	0.69	0.30	99%	28464
1024	Ok	9.24	1.44	0.60	22%	37968



# Transpuesta de matriz



$n \rightarrow$  Es dado como argumento del programa.

$$A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 13 & 14 & 15 & 16 \end{bmatrix} \Rightarrow A^T = \begin{bmatrix} 1 & 5 & 9 & 13 \\ 2 & 6 & 10 & 14 \\ 3 & 7 & 11 & 15 \\ 4 & 8 & 12 & 16 \end{bmatrix}$$

CPU

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
```

```
#define N 4
```

```
void TransponerMatrices( int *h_a, int *h_b, int n )
{
    int id, i;

    for( id = 0 ; id < n ; id++ )
    {
        for( i = 0 ; i < n ; i++ )
        {
            h_b[ id * n + i ] = h_a[ i * n + id ];
        }
    }
}
```



$$A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 13 & 14 & 15 & 16 \end{bmatrix} \Rightarrow A^T = \begin{bmatrix} 1 & 5 & 9 & 13 \\ 2 & 6 & 10 & 14 \\ 3 & 7 & 11 & 15 \\ 4 & 8 & 12 & 16 \end{bmatrix}$$

## GP-GPU

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
```

```
#define N 4
```

```
__global__
void transposerMatrices( int *d_a, int *d_b, int n )
{
    int id, i;

    id = blockIdx.x * blockDim.x + threadIdx.x;

    for( i = 0 ; i < n ; i++ )
    {
        d_b[ id * n + i ] = d_a[ i * n + id ];
    }
}
```



# CPU

```
int main( int argc, char** argv)
{
    int *h_a, *h_b;
    char validacion[10];

    int    n = N; // Valor por defecto

    if ( argc > 1 )
    {
        n = atoi (argv[1]);
        if ( n > 1024 )
        {
            n = 1024;
        }
    }

    size_t memSize = n * n * sizeof( int );

    h_a = (int *) malloc( memSize );
    h_b = (int *) malloc( memSize );

    if ( h_a == NULL || h_b == NULL )
    {
        perror("Memoria insuficiente\n");
        exit(-1);
    }
}
```

# GP-GPU

```
int main( int argc, char** argv)
{
    int *h_a, *h_b;
    int *d_a, *d_b;
    char validacion[10];

    int    n = N; // Valor por defecto

    if ( argc > 1 )
    {
        n = atoi (argv[1]);
        if ( n > 1024 )
        {
            n = 1024;
        }
    }

    size_t memSize = n * n * sizeof( int );

    h_a = (int *) malloc( memSize );
    h_b = (int *) malloc( memSize );

    if ( h_a == NULL || h_b == NULL )
    {
        perror("Memoria insuficiente\n");
        exit(-1);
    }
}
```

**CPU**

```
for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        h_a[ n * i + j ] = n * i + j;
        h_b[ n * i + j ] = 0;
    }
}

TransponerMatrices (h_a, h_b, n );

strcpy(validacion, "Ok");
for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        if ( h_a[ n * i + j ] != h_b[ n * j + i ] )
        {
            strcpy(validacion, "Error");
        }
    }
}
free(h_a);
free(h_b);

printf ( "\t %10d \t\t %s \t ", n, validacion );

return 0;
}
```

# GP-GPU

```
cudaMalloc( (void**) &d_a, memSize );
cudaMalloc( (void**) &d_b, memSize );

if ( d_a == NULL || d_b == NULL )
{
    perror("Memoria insuficiente en la GPU\n");
    exit(-1);
}

for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        h_a[ n * i + j ] = n * i + j;
        h_b[ n * i + j ] = 0;
    }
}


cudaMemcpy( d_a, h_a , memSize, cudaMemcpyHostToDevice );
cudaMemcpy( d_b, h_b , memSize, cudaMemcpyHostToDevice );

transponerMatrices<<< 1 , n >>> (d_a, d_b, n );


cudaMemcpy( h_b, d_b, memSize, cudaMemcpyDeviceToHost );
```

# GP-GPU

```
strcpy(validacion, "Ok");  
for( int i = 0 ; i < n ; i++ )  
{  
    for( int j = 0 ; j < n ; j++ )  
    {  
        if ( h_a[ n * i + j ] != h_b[ n * j + i ] )  
        {  
            strcpy(validacion, "Error");  
        }  
    }  
}  
  
cudaFree(d_a);  
cudaFree(d_b);  
  
free(h_a);  
free(h_b);  
  
printf ( "\t %10d \t\t %s \t ", n, validacion );  
  
return 0;  
}
```




# Producto de matrices - v1



[https://upload.wikimedia.org/wikipedia/commons/thumb/1/11/Matrix\\_multiplication\\_diagram.svg/250px-Matrix\\_multiplication\\_diagram.svg.png](https://upload.wikimedia.org/wikipedia/commons/thumb/1/11/Matrix_multiplication_diagram.svg/250px-Matrix_multiplication_diagram.svg.png)

$n$  -> Es dado como argumento del programa.



# Producto de matrices - v1

CPU


```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
```

```
#define N 4
```

```
void MultiplocarMatrices( int *h_a, int *h_b, int *h_c, int n )
{
    int id, i, j;

    for( id = 0 ; id < n ; id++ )
    {
        for( i = 0 ; i < n ; i++ )
        {
            h_c[ id * n + i ] = 0;
            for ( j = 0 ; j < n ; j++ )
            {
                h_c[ id * n + i ] += ( h_a[ id * n + j ] * h_b[ j * n + id ] );
            }
        }
    }
}
```





# Producto de matrices - v1


**GP-GPU**

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define N 4

__global__
void MultiplicarMatrices(int *d_a, int *d_b, int *d_c, int n )
{
    int id, i, j;

    id = blockIdx.x * blockDim.x + threadIdx.x;
    if ( id < n )
    {
        for( i = 0 ; i < n ; i++ )
        {
            d_c[ id * n + i ] = 0;
            for ( j = 0 ; j < n ; j++ )
            {
                d_c[ id * n + i ] += ( d_a[ id * n + j ] * d_b[ j * n + id ] );
            }
        }
    }
}
```



# Producto de matrices - v1

CPU


```
int main( int argc, char** argv)
{
    int *h_a, *h_b, *h_c;
    char validacion[10];

    int    n = N; // Valor por defecto

    if ( argc > 1 )
    {
        n = atoi (argv[1]);
        if ( n > 1024 )
        {
            n = 1024;
        }
    }

    size_t memSize = n * n * sizeof( int );

    h_a = (int *) malloc( memSize );
    h_b = (int *) malloc( memSize );
    h_c = (int *) malloc( memSize );
```



# Producto de matrices - v1

**GP-GPU**


```
int main( int argc, char** argv)
{
    int *h_a, *h_b, *h_c;
    int *d_a, *d_b, *d_c;
    char validacion[10];

    int    n = N; // Valor por defecto

    if ( argc > 1 )
    {
        n = atoi (argv[1]);
        if ( n > 1024 )
        {
            n = 1024;
        }
    }

    size_t memSize = n * n * sizeof( int );

    h_a = (int *) malloc( memSize );
    h_b = (int *) malloc( memSize );
    h_c = (int *) malloc( memSize );
```




# Producto de matrices - v1

**CPU**

```
if ( h_a == NULL || h_b == NULL || h_c == NULL )
{
    perror("Memoria insuficiente\n");
    exit(-1);
}

for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        h_a[ n * i + j ] = 1;
        h_b[ n * i + j ] = 1;
        h_c[ n * i + j ] = 0;
    }
}

MultiplicarMatrices (h_a, h_b, h_c, n );
```



# Producto de matrices - v1

**GP-GPU**

```
if ( h_a == NULL || h_b == NULL || h_c == NULL )
{
    perror("Memoria insuficiente\n");
    exit(-1);
}
```

```
cudaMalloc( (void**) &d_a, memSize );
cudaMalloc( (void**) &d_b, memSize );
cudaMalloc( (void**) &d_c, memSize );
```


```
if ( d_a == NULL || d_b == NULL || d_c == NULL )
{
    perror("Memoria insuficiente en la GPU\n");
    exit(-1);
}
```

```
for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        h_a[ n * i + j ] = 1;
        h_b[ n * i + j ] = 1;
        h_c[ n * i + j ] = 0;
    }
}
```

```
cudaMemcpy( d_a, h_a , memSize, cudaMemcpyHostToDevice );
cudaMemcpy( d_b, h_b , memSize, cudaMemcpyHostToDevice );

MultiplicarMatrices<<< 1 , n >>> (d_a, d_b, d_c, n );

cudaMemcpy( h_c, d_c, memSize, cudaMemcpyDeviceToHost );
```



# Producto de matrices - v1

CPU


```
strcpy(validacion, "Ok");
for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        if ( h_c[ n * i + j ] != n )
        {
            strcpy(validacion, "Error");
        }
    }
}

free(h_a);
free(h_b);
free(h_c);

printf ( "\t %10d \t\t %s \t ", n, validacion );

return 0;
}
```





# Producto de matrices - v1

**GP-GPU**

```
strcpy(validacion, "Ok");
for( int i = 0 ; i < n ; i++ )
{
    for( int j = 0 ; j < n ; j++ )
    {
        if ( h_c[ n * i + j ] != n )
        {
            strcpy(validacion, "Error");
        }
    }
}


cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

free(h_a);
free(h_b);
free(h_c);

printf ( "\t %10d \t\t %s \t ", n, validacion );

return 0;
}
```






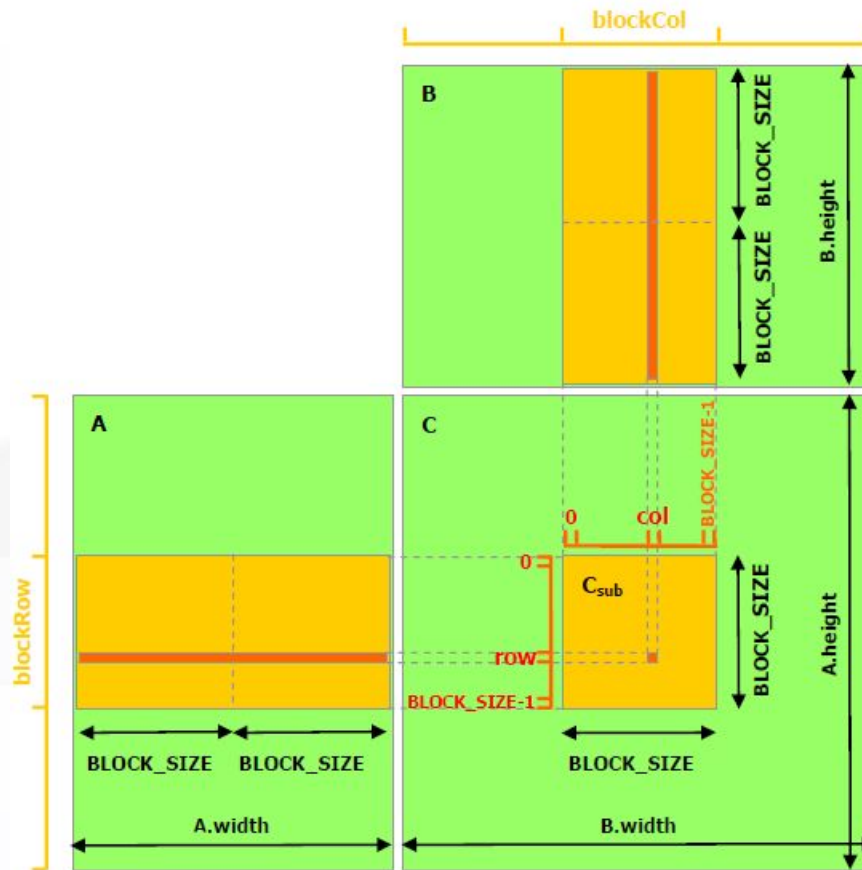
# Producto de matrices - v1

**GP-GPU**


N	Estado	T.Real(sg)	T.Usuario(sg)	T.Kernel(sg)	CPU	Memoria(KB)
2	Ok	0.36	0.01	0.05	18%	25584
4	Ok	0.05	0.00	0.04	94%	25700
8	Ok	0.04	0.00	0.03	100%	25656
16	Ok	0.03	0.00	0.02	97%	25728
32	Ok	0.03	0.00	0.02	91%	25820
64	Ok	0.03	0.00	0.02	91%	23652
128	Ok	0.07	0.03	0.03	97%	27916
256	Ok	0.21	0.13	0.06	97%	26064
512	Ok	0.99	0.69	0.30	99%	28464
1024	Ok	9.24	1.44	0.60	22%	37968



# Producto de matrices - v2




<https://docs.nvidia.com/cuda/cuda-c-programming-guide/graphics/matrix-multiplication-with-shared-memory.png>



# Producto de matrices - v2

**GP-GPU**

```
int MATRIX_SIZE = n;  
int TILE_SIZE = 2;  
  
dim3 dimGrid ( MATRIX_SIZE / TILE_SIZE, MATRIX_SIZE / TILE_SIZE);  
dim3 dimBlock (TILE_SIZE, TILE_SIZE, 1);  
  
MultiplocarMatrices<<< dimGrid , dimBlock >>> (d_a, d_b, d_c, n );
```



# Producto de matrices - v2

**GP-GPU**

```
#define _BLOCK_SIZE_ 2

__global__ void MultiplocarMatrices(int *A, int *B, int *C, int n)
{

    const uint wA = n;
    const uint wB = n;


    const uint bx = blockIdx.x;
    const uint by = blockIdx.y;

    const uint tx = threadIdx.x;
    const uint ty = threadIdx.y;

    const uint aBegin = wA * _BLOCK_SIZE_ * by;
    const uint aEnd   = aBegin + wA - 1;
    const uint aStep  = _BLOCK_SIZE_;

    const uint bBegin = _BLOCK_SIZE_ * bx;
    const uint bStep  = _BLOCK_SIZE_ * wB;

    float Csub = 0;
```



# Producto de matrices - v2

**GP-GPU**

```
for (int a = aBegin, b = bBegin; a <= aEnd; a += aStep, b += bStep)
{
    __shared__ float As[_BLOCK_SIZE][_BLOCK_SIZE];
    __shared__ float Bs[_BLOCK_SIZE][_BLOCK_SIZE];


    As[ty][tx] = A[a + wA * ty + tx];
    Bs[ty][tx] = B[b + wB * ty + tx];

    __syncthreads();

    for (int k = 0; k < _BLOCK_SIZE; ++k)
        Csub += As[ty][k] * Bs[k][tx];


    __syncthreads();
}

const uint c = wB * _BLOCK_SIZE * by + _BLOCK_SIZE * bx;
C[c + wB * ty + tx] = Csub;
}
```



# Producto de matrices - v2

N	Estado	T.Real(sg)	T.Usuario(sg)	T.Kernel(sg)	CPU	Memoria(KB)
2	Ok	0.03	0.00	0.02	71%	26860
4	Ok	0.04	0.01	0.02	95%	26908
8	Ok	0.03	0.01	0.02	92%	24912
16	Ok	0.03	0.00	0.02	94%	24760
32	Ok	0.03	0.00	0.02	93%	24732
64	Ok	0.03	0.00	0.02	90%	24720
128	Ok	0.03	0.02	0.01	92%	25056
256	Ok	0.12	0.06	0.05	97%	25288
512	Ok	0.80	0.56	0.24	99%	27516
1024	Ok	7.07	5.05	2.01	99%	36944



# Gracias